



6 - Deploy Video Processing Service



10:25

14:28



Mark Lesson Complete



Full Stack Development

19 / 22

Intro

0

Demo and
8 min FREE
Architecture

1

F 4 min L FREE

Video Processing Service

Initialize

Host Video Processing Service on Google Cloud Run

We will now deploy our service to Google **Cloud Run**, which is a fully-managed serverless container platform. That means we don't need to worry about managing servers or scaling our service.

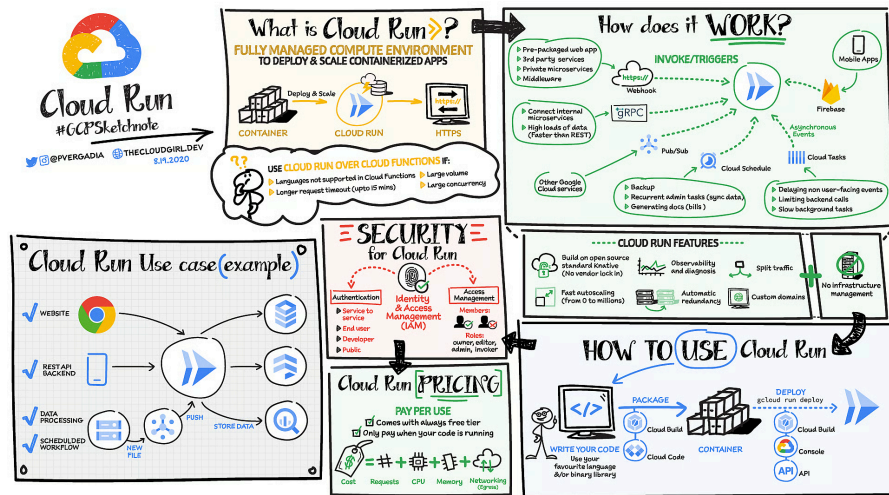
2 Video 11 min FREE
Processing
Service

3 Process 13 min FREE
Locally

4 Video 15 min
Processing
Service

5 Convert
Videos
Hosted
on 24 min
Google
Cloud
Storage

Google Cloud



There are a few steps involved:

1. Create a Google Cloud Project and enable billing (in our case, we will create a Firebase project).
2. Install `gcloud` and `gsutil` CLI tools
3. Upload our Docker image to Google Artifact Registry
4. Deploy our Docker image to Google Cloud Run

1. Create a Firebase Project

We will create a Firebase project, which will also create a Google Cloud project for us.

You can do this via the Firebase console:

<https://console.firebase.google.com/>

Take note of the `project id` since we will need it throughout the course.

This will create a Google Cloud project for us with the same `project id`, which we can see in the Google Cloud console:

<https://console.cloud.google.com/>

Google Cloud offers a free tier, as well as \$300 of free credit. Which should be more than enough for this course.

<https://cloud.google.com/free/docs/free-cloud-features>

2. Install `gcloud` and `gsutil` CLI tools

Install `gcloud` and `gsutil` CLI tools. Then initialize `gcloud` and authenticate with your project.

The steps will depend on your operating system, so it's best to refer to the official docs.

<https://cloud.google.com/sdk/docs/install>

Make sure to authentication with your account and set your project:

```
gcloud auth login # Copy the output url and  
paste it into your browser
```

```
gcloud config set project <PROJECT_ID>
```

3. Upload our Docker image to Google Artifact Registry

Enable artifact registry

```
gcloud services enable  
artifactregistry.googleapis.com
```

(Optional) Update gcloud components

```
gcloud components update
```

Create an Artifact Registry repository:

```
gcloud artifacts repositories create video-  
processing-repo \  
  --repository-format=docker \  
  --location=us-central1 \  
  --description="Docker repository for video  
processing service"
```

For some reason I had an issue with this command, so I created the repo via the console:

<https://console.cloud.google.com/artifacts>

Then, rebuild your Docker image. With this naming scheme, docker knows where to push the image and which project.

```
docker build -t us-central1-  
docker.pkg.dev/<PROJECT_ID>/video-processing-  
repo/video-processing-service .
```

Important: If you are using mac, add `--platform linux/amd64` to the above command.

You may still need to configure Docker to use gcloud as the credential helper:

```
gcloud auth configure-docker us-central1-  
docker.pkg.dev
```

Then, push the Docker image to Google Artifact Registry:

```
docker push us-central1-  
docker.pkg.dev/<PROJECT_ID>/video-processing-  
repo/video-processing-service
```

4. Deploy our Docker image to Google Cloud Run

Lastly, deploy to Cloud Run.

```
# Enable cloud run  
gcloud services enable run.googleapis.com  
  
# Deploy container to cloud run  
gcloud run deploy video-processing-service --  
image us-central1-  
docker.pkg.dev/PROJECT_ID/video-processing-  
repo/video-processing-service \  
  --region=us-central1 \  
  --platform managed \  
  --timeout=3600 \  
  --memory=2Gi \  
  --cpu=1 \  
  --min-instances=0 \  
  --max-instances=1 \  
  --ingress=internal
```

We are setting the ingress to internal so that only GCP internal services can access it. This is so that our Pub/Sub service can invoke it, but not the outside world.

In addition, it might be good to use the `--no-allow-unauthenticated` flag to limit other internal services. Meaning we can set it up so that only our Pub/Sub service can invoke it and not our other GCP internal services. But this would be cumbersome to setup, so we will skip it for now.